

DS542 Midterm Report

Yuchen Li

Spring 2025

1 Part 1 – SimpleCNN

For the baseline, I manually defined a SimpleCNN model with:

- Three convolutional blocks with **MaxPooling** after each block
- No data augmentation
- Adam optimizer with **StepLR** scheduler
- Learning rate set to 0.1
- Batch size of 8

The result was suboptimal. Within 5 epochs, training accuracy was only 0.995, and validation accuracy was just 1.08%.

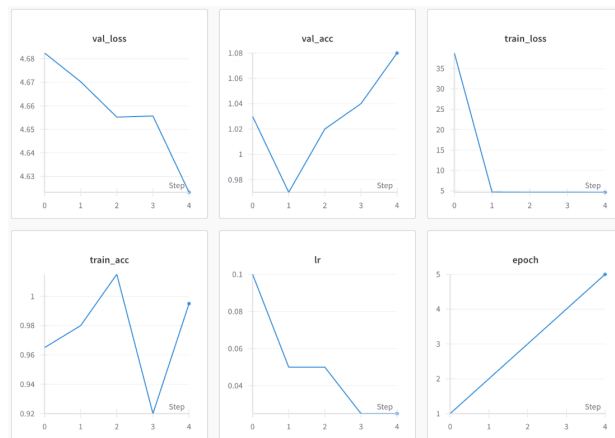


Figure 1: Evaluation Metrics for SimpleCNN (Before Tunning)

To improve model performance, I made several key changes:

- Built a deeper network, added Batch Normalization for faster convergence and generalization, Dropout layers to prevent co-adaptation, and used Global Average Pooling instead of flattening
- Applied data augmentation to introduce randomness (`RandomHorizontalFlip()` and `RandomCrop(32, padding=4)`)

- Incorporated L2 regularization via `weight_decay` in the `Adam` optimizer. Increased batch size to 64 and lowered the learning rate.

Under these modifications, the `SimpleCNN` reached 36.02% training accuracy and 38.83% validation accuracy within 5 epochs.

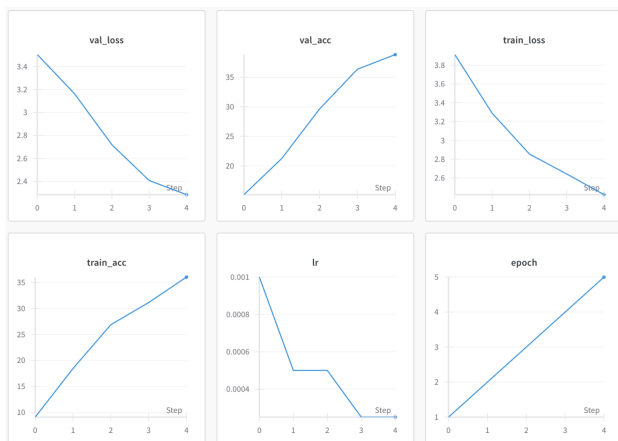


Figure 2: Evaluation Metrics for SimpleCNN (After Tuning)

2 Part 2 – ResNet18

In the second part, I transitioned to a more sophisticated architecture: `ResNet18` from `torchvision.models`, modifying the final fully connected layer to output 100 classes. This choice was motivated by its lightweight nature and residual connections that aid gradient flow.

Without further changes, switching from `SimpleCNN` to `ResNet18` resulted in training accuracy of 37.55% and validation accuracy of 36.26%.

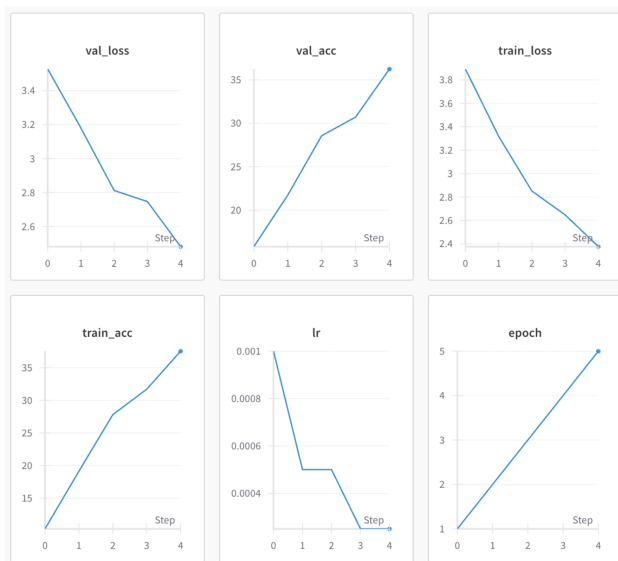


Figure 3: Evaluation Metrics for Part 2 (Without Further Tuning)

Improvements Attempted

First Try:

1. Replaced Adam optimizer with SGD + Momentum for better generalization and gradient smoothing.
2. Used CosineAnnealingLR instead of StepLR for smoother LR decay.

```
1 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9,
2   weight_decay=1e-4)
3 scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=
4   CONFIG["epochs"])
```

3. Enhanced data augmentation with ColorJitter and RandomRotation.
4. Applied label smoothing in the loss function.

```
1 criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
```

However, these changes did not yield better results within 5 epochs. The model achieved only 14.20% training accuracy and 20.87% validation accuracy.

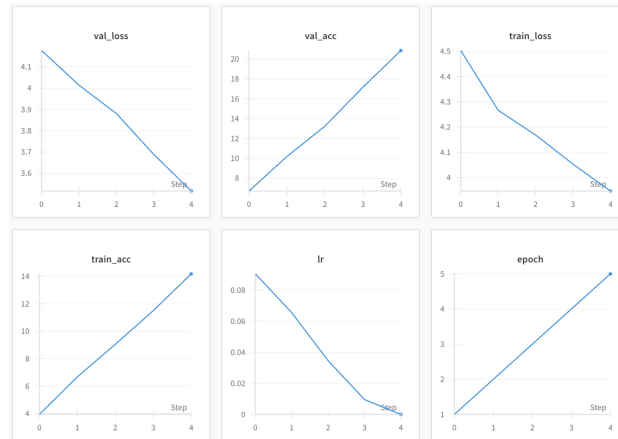


Figure 4: Evaluation Metrics for Part 2 (1st Try)

Second Try:

Introduced **Mixup** for advanced data augmentation. Since ResNet18 is deep and prone to overfitting, Mixup helped reduce overfitting by smoothing decision boundaries and improving generalization. The technique was implemented as:

```
1 def mixup_data(x, y, alpha=1.0):
2     if alpha > 0:
3         lam = np.random.beta(alpha, alpha)
4     else:
5         lam = 1.0
6     batch_size = x.size(0)
7     index = torch.randperm(batch_size).to(x.device)
```

```

8
9     mixed_x = lam * x + (1 - lam) * x[index, :]
10    y_a, y_b = y, y[index]
11    return mixed_x, y_a, y_b, lam

```

This yielded:

Train accuracy: 25.91%, Validation accuracy: 25.06%.

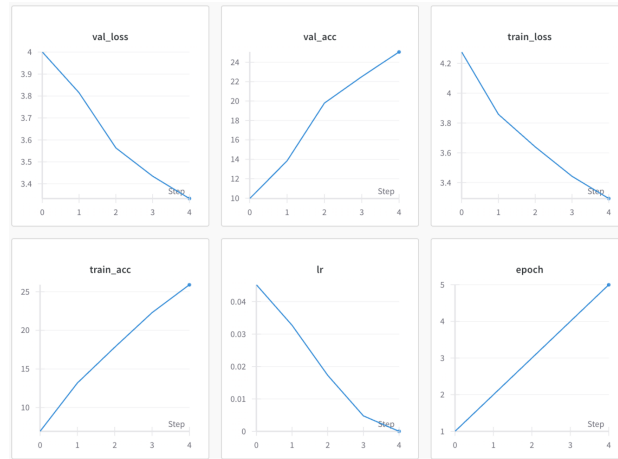


Figure 5: Evaluation Metrics for Part 2 (2nd Try)

Third Try:

Replaced `CosineAnnealingLR` with `OneCycleLR`, which is often effective with SGD, for it helps models converge faster and generalize better.

```

1     scheduler = optim.lr_scheduler.OneCycleLR(
2         optimizer,
3         max_lr=CONFIG["learning_rate"],
4         steps_per_epoch=len(trainloader),
5         epochs=CONFIG["epochs"],
6         pct_start=0.3,
7         anneal_strategy='cos',
8         div_factor=25.0,
9         final_div_factor=1e4
10    )

```

However, the result was worse than before, achieving only 20.73% training accuracy and 21.32% validation accuracy.

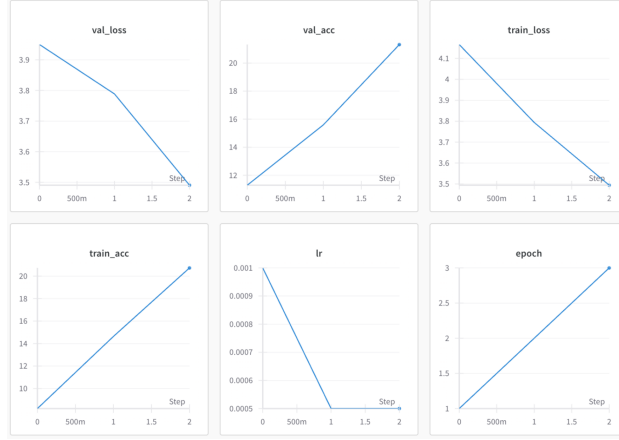


Figure 6: Evaluation Metrics for Part 2 (3rd Try)

Final Decision:

- Reverted to Adam optimizer with StepLR
- Removed Mixup

This configuration achieved the best result in 5 epochs:

Train accuracy: 40.21%, Validation accuracy: 41.3%

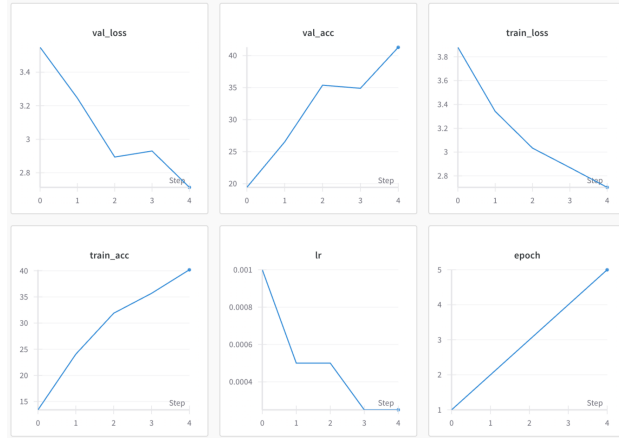


Figure 7: Evaluation Metrics for Part 2 (Final Version)

3 Part 3 – MobileNetV2

For the final part, I leveraged transfer learning by using MobileNetV2 with `pretrained=True`, resizing CIFAR-100 images to 224×224 and increasing batch size to 128.

Reasons for choosing MobileNetV2:

- Efficient, fast, and memory-friendly
- Uses depthwise separable convolutions to reduce computational cost

- Well-suited for training on SCC with limited GPU resources

With 50 epochs, the final result was:

Train accuracy: 99.95%, Validation accuracy: 78.25%

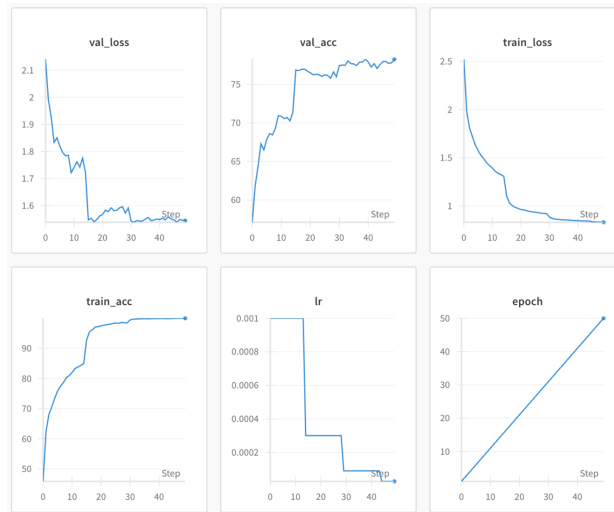


Figure 8: Evaluation Metrics for MobileNetV2

Challenges and Debugging

Initially, the predictions in the OOD submission CSV were all the same. After debugging, I discovered the problem stemmed from how inputs were unpacked in the TensorDataset and normalized in batch. With the original code:

```
1 with torch.no_grad():
2     for inputs in tqdm(dataloader, desc=f"Evaluating {distortion_name}
3         (Severity {severity})", leave=False):
4         inputs = inputs[0]
5         inputs = normalize(inputs)
6         inputs = inputs.to(device)
```

Since each inputs from TensorDataset is a tuple, if it repeats the process in a loop, it might return the same tensor reference repeatedly. Also, the code normalizes the whole batch tensor without looping. So I fix that into:

```
1 normalize = transforms.Normalize((0.5071, 0.4867, 0.4408), (0.2675,
2     0.2565, 0.2761))
3 resize = transforms.Resize((224, 224))
4 images = torch.stack([normalize(resize(img)) for img in images])
5 ...
6 with torch.no_grad():
7     for batch in tqdm(dataloader, desc=f"Evaluating {distortion_name}
8         (Severity {severity})", leave=False):
9         inputs = batch[0].to(device)
10        outputs = model(inputs)
```

```
9         _, predicted = outputs.max(1)
10        predictions.extend(predicted.cpu().numpy())
```

Once corrected, the submission results became reasonable.

4 Future Improvements

If I had more time, I would explore:

- EfficientNet B0 with transfer learning
- CutMix or Stochastic Depth
- Learning rate finder for auto-tuning
- Extended training schedules (100+ epochs) with checkpointing
- Early stopping and ensemble modeling